

Spark Tuning Checklist

1. Set `spark.sql.shuffle.partitions` to 2-4x the number of cores.
2. Enable AQE (`spark.sql.adaptive.enabled=true`).
3. Use Parquet with Snappy compression for large datasets.
4. Avoid `collect()` on large DataFrames; use `show()` or `write` instead.